

# Classification Algorithms – Regression & Knn

# Outline

- **Linear Models (Regression)**
- Instance-based (Nearest-neighbor)



# Linear models

- Work most naturally with numeric attributes
- Standard technique for numeric prediction: linear regression
  - Outcome is linear combination of attributes

$$x = w_0 + w_1 a_1 + w_2 a_2 + \dots + w_k a_k$$

- Weights are calculated from the training data
- Predicted value for first training instance  $\mathbf{a}^{(1)}$

$$w_0 a_0^{(1)} + w_1 a_1^{(1)} + w_2 a_2^{(1)} + \dots + w_k a_k^{(1)} = \sum_{j=0}^k w_j a_j^{(1)}$$

# Minimizing the squared error

- Choose  $k + 1$  coefficients to minimize the squared error on the training data

- Squared error: 
$$\sum_{i=1}^n \left( x^{(i)} - \sum_{j=0}^k w_j a_j^{(i)} \right)^2$$

- Derive coefficients using standard matrix operations
- Can be done if there are more instances than attributes (roughly speaking)
- Minimizing the *absolute error* is more difficult

# Regression for Classification

- *Any* regression technique can be used for classification
  - Training: perform a regression for each class, setting the output to 1 for training instances that belong to class, and 0 for those that don't
  - Prediction: predict class corresponding to model with largest output value (*membership value*)
- For linear regression this is known as *multi-response linear regression*

# \*Theoretical justification

*Observed target value (either 0 or 1)*

*Model*

*Instance*

*The scheme minimizes this*

*True class probability*

$$\begin{aligned}
 & E_y \{ (f(X) - Y)^2 \mid X = x \} \\
 &= E_y \{ (f(X) - P(Y = 1 \mid X = x) + P(Y = 1 \mid X = x) - Y)^2 \mid X = x \} \\
 &= (f(x) - P(Y = 1 \mid X = x))^2 + 2 \times (f(x) - P(Y = 1 \mid X = x)) \times \\
 &\quad E_y \{ P(Y = 1 \mid X = x) - Y \mid X = x \} + E_y \{ (P(Y = 1 \mid X = x) - Y)^2 \mid X = x \} \\
 &= (f(x) - P(Y = 1 \mid X = x))^2 + 2 \times (f(x) - P(Y = 1 \mid X = x)) \times \\
 &\quad (P(Y = 1 \mid X = x) - E_y \{ Y \mid X = x \}) + E_y \{ (P(Y = 1 \mid X = x) - Y)^2 \mid X = x \} \\
 &= \underbrace{(f(x) - P(Y = 1 \mid X = x))^2}_{\text{We want to minimize this}} + \underbrace{E_y \{ (P(Y = 1 \mid X = x) - Y)^2 \mid X = x \}}_{\text{Constant}}
 \end{aligned}$$

# \*Pairwise regression

- Another way of using regression for classification:
  - A regression function for every *pair* of classes, using only instances from these two classes
  - Assign output of +1 to one member of the pair, -1 to the other
- Prediction is done by voting
  - Class that receives most votes is predicted
  - Alternative: "don't know" if there is no agreement
- More likely to be accurate but more expensive

# Logistic regression

- Problem: some assumptions violated when linear regression is applied to classification problems
- *Logistic* regression: alternative to linear regression
  - Designed for classification problems
  - Tries to estimate class probabilities directly
    - Does this using the *maximum likelihood* method
  - Uses this linear model:

$$\log\left(\frac{P}{1-P}\right) = w_0a_0 + w_1a_1 + w_2a_2 + \dots + w_k a_k$$

*P* = Class probability



# Discussion of linear models

- Not appropriate if data exhibits non-linear dependencies
- But: can serve as building blocks for more complex schemes (i.e. model trees)
- Example: multi-response linear regression defines a *hyperplane* for any two given classes:

$$(w_0^{(1)} - w_0^{(2)})a_0 + (w_1^{(1)} - w_1^{(2)})a_1 + (w_2^{(1)} - w_2^{(2)})a_2 + \dots + (w_k^{(1)} - w_k^{(2)})a_k > 0$$

# Comments on basic methods

- Minsky and Papert (1969) showed that linear classifiers have limitations, e.g. can't learn XOR
  - But: combinations of them can (→ Neural Nets)

# Outline

- Linear Models (Regression)
- **Instance-based (Nearest-neighbor)**



# Instance-based representation

- Simplest form of learning: *vote learning*
  - Training instances are searched for instance that most closely resembles new instance
  - The instances themselves represent the knowledge
  - Also called ***instance-based*** learning
- Similarity function defines what's "learned"
- Instance-based learning is ***lazy*** learning
- Methods:
  - *nearest-neighbor*
  - *k-nearest-neighbor*
  - ...

# The distance function

- Simplest case: one numeric attribute
  - Distance is the difference between the two attribute values involved (or a function thereof)
- Several numeric attributes: normally, Euclidean distance is used and attributes are normalized
- Nominal attributes: distance is set to 1 if values are different, 0 if they are equal
- Are all attributes equally important?
  - Weighting the attributes might be necessary

# Instance-based learning

- Distance function defines what's learned
- Most instance-based schemes use *Euclidean distance*:

$$\sqrt{(a_1^{(1)} - a_1^{(2)})^2 + (a_2^{(1)} - a_2^{(2)})^2 + \dots + (a_k^{(1)} - a_k^{(2)})^2}$$

$\mathbf{a}^{(1)}$  and  $\mathbf{a}^{(2)}$ : two instances with  $k$  attributes

- Taking the square root is not required when comparing distances
- Other popular metric: *city-block (Manhattan) metric*
  - Adds differences without squaring them

# Normalization and other issues

- Different attributes are measured on different scales  $\Rightarrow$  need to be *normalized*:

$$a_i = \frac{v_i - \min v_i}{\max v_i - \min v_i} \quad \text{or} \quad a_i = \frac{v_i - \text{Avg}(v_i)}{\text{StDev}(v_i)}$$

$v_i$ : the actual value of attribute  $i$

- Nominal attributes: distance either 0 or 1
- Common policy for missing values: assumed to be maximally distant (given normalized attributes)

# Discussion of 1-NN

- Often very accurate
- ... but slow:
  - simple version scans entire training data to derive a prediction
- Assumes all attributes are equally important
  - Remedy: attribute selection or weights
- Possible remedies against noisy instances:
  - Take a majority vote over the  $k$  nearest neighbors
  - Removing noisy instances from dataset (difficult!)
- Statisticians have used  $k$ -NN since early 1950s
  - If  $n \rightarrow \infty$  and  $k/n \rightarrow 0$ , error approaches minimum



# Summary

- Simple methods frequently work well
  - robust against noise, errors
- Advanced methods, if properly used, can improve on simple methods
- No method is universally best